



NVIDIA Virtual GPU Software What's New

Release 20.0-20.2

NVIDIA Corporation

Sep 30, 2026

Contents

1	Updates by Release	1
1.1	Updates in Release 20.2	1
1.1.1	New Features in Release 20.2	1
1.1.2	Hardware and Software Support Introduced in Release 20.2	1
1.1.3	Features Deprecated in Release 20.2	2
1.1.4	Feature Support Withdrawn in Release 20.2	2
1.2	Updates in Release 20.1	2
1.2.1	New Features in Release 20.1	2
1.2.2	Hardware and Software Support Introduced in Release 20.1	2
1.2.3	Feature Support Withdrawn in Release 20.1	3
1.3	Updates in Release 20.0	3
1.3.1	New Features in Release 20.0	3
1.3.2	Newly Supported Hardware and Software in Release 20.0	3
1.3.3	Feature Support Withdrawn in Release 20.0	4
2	Resolved Issues	5
2.1	Issues Resolved in Release 20.2	5
2.2	Issues Resolved in Release 20.1	5
2.3	Issues Resolved in Release 20.0	5
3	Notices	7

Chapter 1. Updates by Release

Updates for each release in this release family of NVIDIA vGPU software may include new features, introduction of hardware and software support, and withdrawal of hardware and software support.

1.1. Updates in Release 20.2

1.1.1. New Features in Release 20.2

- ▶ Security updates

When these updates are disclosed, full details are provided in a new public NVIDIA GPU Display Driver security bulletin, which is listed on the [NVIDIA Product Security](#) page. To be notified of the availability of this and future bulletins, use the link on the [NVIDIA Product Security](#) page to subscribe to receive email notifications of security bulletin updates.

- ▶ Support for the assignment of multiple MIG-backed vGPUs to the same VM on VCF 9.1 or a later release
- ▶ Miscellaneous bug fixes

1.1.2. Hardware and Software Support Introduced in Release 20.2

- ▶ Newly supported graphics cards:

- ▶ NVIDIA RTX PRO 5000 Blackwell Workstation Edition 72 GB on Red Hat Enterprise Linux with KVM 9.6 only

- ▶ Newly supported hypervisor software:

- ▶ XenServer 9

The Virtual GPU Manager for this release is supplied **only** as a Supplemental Pack. No RPM package is supplied for this release.

- ▶ Red Hat Enterprise Linux with KVM 10.2
 - ▶ Red Hat Enterprise Linux with KVM 9.8

- ▶ Newly supported guest OS releases:

- ▶ Red Hat Enterprise Linux 10.2
 - ▶ Red Hat Enterprise Linux 9.8

- ▶ Newly supported remoting solutions:
 - ▶ Citrix Virtual Apps and Desktops version 7 2603

1.1.3. Features Deprecated in Release 20.2

The following table lists features the use of which is discouraged starting with this release of NVIDIA vGPU software. Although the features remain available, for enhanced security, you should use the preferred alternative listed in the table.

Deprecated Feature	Preferred native	Alter-	Additional Information
RPM package installation of the Virtual GPU Manager for XenServer 8.4	Supplemental Pack installation		Installing or Updating the Supplemental Pack for XenServer

1.1.4. Feature Support Withdrawn in Release 20.2

- ▶ Hypervisor software no longer supported:
 - ▶ Microsoft Azure Local 2411
 - ▶ Red Hat Enterprise Linux with KVM 9.7
- ▶ Guest OS releases no longer supported:
 - ▶ Red Hat Enterprise Linux 9.7

1.2. Updates in Release 20.1

1.2.1. New Features in Release 20.1

- ▶ Security updates

When these updates are disclosed, full details are provided in a new public NVIDIA GPU Display Driver security bulletin, which is listed on the [NVIDIA Product Security](#) page. To be notified of the availability of this and future bulletins, use the link on the [NVIDIA Product Security](#) page to subscribe to receive email notifications of security bulletin updates.

- ▶ Miscellaneous bug fixes as explained in [Issues Resolved in Release 20.1](#)

1.2.2. Hardware and Software Support Introduced in Release 20.1

- ▶ Newly supported hypervisor software:
 - ▶ Microsoft Azure Local 24H2 (12.2603.1002.500)
 - ▶ Ubuntu 26.04 LTS
- ▶ Newly supported guest OS releases:

- ▶ SUSE Linux Enterprise Server 16.0
- ▶ SUSE Linux Enterprise Server 15 SP7
- ▶ SUSE Linux Enterprise Server 15 SP6
- ▶ Ubuntu 26.04 LTS
- ▶ Newly supported remoting solutions:
 - ▶ Omnissa Horizon 8 2603

1.2.3. Feature Support Withdrawn in Release 20.1

- ▶ Guest OS releases no longer supported:
 - ▶ SUSE Linux Enterprise Server 15 SP2

1.3. Updates in Release 20.0

1.3.1. New Features in Release 20.0

- ▶ Support for the fixed share scheduler for time-sliced vGPUs with different amounts of frame buffer on the same physical GPU
- ▶ Support for the Wayland display server protocol
- ▶ Support for the following features on VCF 9.1:
 - ▶ Multi-Instance GPU (MIG)-backed vGPUs for graphics on GPUs that support MIG
 - ▶ Time-sliced, MIG-backed vGPUs within a GPU instance on a MIG-enabled GPU
- ▶ Miscellaneous bug fixes as explained in *Issues Resolved in Release 20.0*

1.3.2. Newly Supported Hardware and Software in Release 20.0

- ▶ Newly supported graphics cards:
 - ▶ NVIDIA RTX PRO 6000 Blackwell Server Edition liquid cooled
 - ▶ NVIDIA RTX PRO 4500 Blackwell Server Edition
- ▶ Newly supported hypervisor software:
 - ▶ VMware Cloud Foundation (VCF) 9.1
- ▶ Newly supported generic Linux with KVM platforms:
 - ▶ Verge.io VergeOS
- ▶ Newly supported guest OS releases:
 - ▶ Red Hat Enterprise Linux 10.1 and 10.0 on VCF since VCF 9.1

1.3.3. Feature Support Withdrawn in Release 20.0

- ▶ Disabling strict round robin policy is no longer supported.
- ▶ The following graphics cards are no longer supported:
 - ▶ Tesla M10
 - ▶ Tesla V100 SXM2
 - ▶ Tesla V100 SXM2 32GB
 - ▶ Tesla V100 PCIe
 - ▶ Tesla V100 PCIe 32GB
 - ▶ Tesla V100S PCIe 32GB
 - ▶ Tesla V100 FHHL
 - ▶ Quadro RTX 6000
 - ▶ Quadro RTX 6000 passive
 - ▶ Quadro RTX 8000
 - ▶ Quadro RTX 8000 passive

Chapter 2. Resolved Issues

Only resolved issues that have been previously noted as known issues or had a noticeable user impact are listed. The summary and description for each resolved issue indicate the effect of the issue on NVIDIA vGPU software **before the issue was resolved**.

2.1. Issues Resolved in Release 20.2

No resolved issues are reported in this release for all supported hypervisors.

2.2. Issues Resolved in Release 20.1

No resolved issues are reported in this release for all supported hypervisors.

2.3. Issues Resolved in Release 20.0

No resolved issues are reported in this release for all supported hypervisors.

Chapter 3. Notices

Notice

This document is provided for information purposes only and shall not be regarded as a warranty of a certain functionality, condition, or quality of a product. NVIDIA Corporation (“NVIDIA”) makes no representations or warranties, expressed or implied, as to the accuracy or completeness of the information contained in this document and assumes no responsibility for any errors contained herein. NVIDIA shall have no liability for the consequences or use of such information or for any infringement of patents or other rights of third parties that may result from its use. This document is not a commitment to develop, release, or deliver any Material (defined below), code, or functionality.

NVIDIA reserves the right to make corrections, modifications, enhancements, improvements, and any other changes to this document, at any time without notice.

Customer should obtain the latest relevant information before placing orders and should verify that such information is current and complete.

NVIDIA products are sold subject to the NVIDIA standard terms and conditions of sale supplied at the time of order acknowledgement, unless otherwise agreed in an individual sales agreement signed by authorized representatives of NVIDIA and customer (“Terms of Sale”). NVIDIA hereby expressly objects to applying any customer general terms and conditions with regards to the purchase of the NVIDIA product referenced in this document. No contractual obligations are formed either directly or indirectly by this document.

NVIDIA products are not designed, authorized, or warranted to be suitable for use in medical, military, aircraft, space, or life support equipment, nor in applications where failure or malfunction of the NVIDIA product can reasonably be expected to result in personal injury, death, or property or environmental damage. NVIDIA accepts no liability for inclusion and/or use of NVIDIA products in such equipment or applications and therefore such inclusion and/or use is at customer’s own risk.

NVIDIA makes no representation or warranty that products based on this document will be suitable for any specified use. Testing of all parameters of each product is not necessarily performed by NVIDIA. It is customer’s sole responsibility to evaluate and determine the applicability of any information contained in this document, ensure the product is suitable and fit for the application planned by customer, and perform the necessary testing for the application in order to avoid a default of the application or the product. Weaknesses in customer’s product designs may affect the quality and reliability of the NVIDIA product and may result in additional or different conditions and/or requirements beyond those contained in this document. NVIDIA accepts no liability related to any default, damage, costs, or problem which may be based on or attributable to: (i) the use of the NVIDIA product in any manner that is contrary to this document or (ii) customer product designs.

No license, either expressed or implied, is granted under any NVIDIA patent right, copyright, or other NVIDIA intellectual property right under this document. Information published by NVIDIA regarding third-party products or services does not constitute a license from NVIDIA to use such products or services or a warranty or endorsement thereof. Use of such information may require a license from a

third party under the patents or other intellectual property rights of the third party, or a license from NVIDIA under the patents or other intellectual property rights of NVIDIA.

Reproduction of information in this document is permissible only if approved in advance by NVIDIA in writing, reproduced without alteration and in full compliance with all applicable export laws and regulations, and accompanied by all associated conditions, limitations, and notices.

THIS DOCUMENT AND ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE. TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL NVIDIA BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF NVIDIA HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES. Notwithstanding any damages that customer might incur for any reason whatsoever, NVIDIA's aggregate and cumulative liability towards customer for the products described herein shall be limited in accordance with the Terms of Sale for the product.

VESA DisplayPort

DisplayPort and DisplayPort Compliance Logo, DisplayPort Compliance Logo for Dual-mode Sources, and DisplayPort Compliance Logo for Active Cables are trademarks owned by the Video Electronics Standards Association in the United States and other countries.

HDMI

HDMI, the HDMI logo, and High-Definition Multimedia Interface are trademarks or registered trademarks of HDMI Licensing LLC.

OpenCL

OpenCL is a trademark of Apple Inc. used under license to the Khronos Group Inc.

Trademarks

NVIDIA, the NVIDIA logo, NVIDIA GRID, NVIDIA GRID vGPU, NVIDIA Maxwell, NVIDIA Pascal, NVIDIA Turing, NVIDIA Volta, GPUDirect, Quadro, and Tesla are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

©2013-2026, NVIDIA Corporation